

Lecture 7: 26 April, 2021


Madhavan Mukund

<https://www.cmi.ac.in/~madhavan>

Data Mining and Machine Learning
April–July 2021


Linear regression

- Find the line that “fits” the data best
 - Normal equation
 - Gradient descent
- Linear: each parameter's contribution is independent
- Input $x : (x_1, x_2, \dots, x_k)$
- $y = \theta_0 + \theta_1 x_1 + \dots + \theta_k x_k$



The non-linear case


- What if the relationship is not linear?
- Here the best possible explanation seems to be a quadratic



The non-linear case

- What if the relationship is not linear?
- Here the best possible explanation seems to be a quadratic
- Non-linear : cross dependencies
- Input $x_i : (x_{i_1}, x_{i_2})$
- Quadratic dependencies:

$$y = \theta_0 + \theta_1 x_{i_1} + \theta_2 x_{i_2} + \theta_{11} x_{i_1}^2 + \theta_{22} x_{i_2}^2 + \theta_{12} x_{i_1} x_{i_2}$$




The non-linear case

- Recall how we fit a line

$$\begin{bmatrix} 1 & x_i \end{bmatrix} \begin{bmatrix} \theta_0 \\ \theta_1 \end{bmatrix}$$

- For quadratic, add new coefficients and expand parameters

$$\begin{bmatrix} 1 & x_i & x_i^2 \end{bmatrix} \begin{bmatrix} \theta_0 \\ \theta_1 \\ \theta_2 \end{bmatrix}$$




The non-linear case

- Input (x_{i1}, x_{i2})
- For the general quadratic case, we are adding new derived “features”

$$x_{i3} = x_{i1}^2$$

$$x_{i4} = x_{i2}^2$$


$$x_{i5} = x_{i1} x_{i2}$$



The non-linear case

- Original input matrix

$$\begin{bmatrix} 1 & x_{1_1} & x_{1_2} \\ 1 & x_{2_1} & x_{2_2} \\ & \dots & \\ 1 & x_{i_1} & x_{i_2} \\ & \dots & \\ 1 & x_{n_1} & x_{n_2} \end{bmatrix}$$




The non-linear case

- Expanded input matrix

$$\begin{bmatrix} 1 & x_{1_1} & x_{1_2} & x_{1_1}^2 & x_{1_2}^2 & x_{1_1}x_{1_2} \\ 1 & x_{2_1} & x_{2_2} & x_{2_1}^2 & x_{2_2}^2 & x_{2_1}x_{2_2} \\ \dots & \dots & \dots & \dots & \dots & \dots \\ 1 & x_{i_1} & x_{i_2} & x_{i_1}^2 & x_{i_2}^2 & x_{i_1}x_{i_2} \\ \dots & \dots & \dots & \dots & \dots & \dots \\ 1 & x_{n_1} & x_{n_2} & x_{n_1}^2 & x_{n_2}^2 & x_{n_1}x_{n_2} \end{bmatrix}$$

- New columns are computed and filled in from original inputs



Exponential parameter blow-up

■ Cubic derived features

$$x_{i_1}^3, x_{i_2}^3, x_{i_3}^3,$$

$$x_{i_1}^2 x_{i_2}, x_{i_1}^2 x_{i_3},$$

$$x_{i_2}^2 x_{i_1}, x_{i_2}^2 x_{i_3},$$


$$x_{i_3}^2 x_{i_1}, x_{i_3}^2 x_{i_2},$$

$$x_{i_1} x_{i_2} x_{i_3},$$

$$x_{i_1}^2, x_{i_2}^2, x_{i_3}^2,$$


$$x_{i_1} x_{i_2}, x_{i_1} x_{i_3}, x_{i_2} x_{i_3},$$

$$x_{i_1}, x_{i_2}, x_{i_3}.$$




Higher degree polynomials

- How complex a polynomial should we try?
- Aim for degree that minimizes SSE
- As degree increases, features explode exponentially



Overfitting


- Need to be careful about adding higher degree terms
- For n training points, can always fit polynomial of degree $(n - 1)$ exactly
- However, such a curve would not generalize well to new data points
- **Overfitting** — model fits training data well, performs poorly on unseen data



Regularization

- Need to trade off SSE against curve complexity
- So far, the only cost has been SSE
- Add a cost related to parameters $(\theta_0, \theta_1, \dots, \theta_k)$
- Minimize, for instance

$$\frac{1}{2} \sum_{i=1}^n (z_i - y_i)^2 + \sum_{j=1}^k \theta_j^2$$



Regularization


$$\frac{1}{2} \sum_{i=1}^n (z_i - y_i)^2 + \sum_{j=1}^k \theta_j^2$$

- Second term penalizes curve complexity
- Variations on regularization

- Ridge regression: $\sum_{j=1}^k \theta_j^2$


- LASSO regression: $\sum_{j=1}^k |\theta_j|$

- Elastic net regression: $\sum_{j=1}^k \lambda_1 |\theta_j| + \lambda_2 \theta_j^2$




The non-polynomial case

- Percentage of urban population as a function of per capita GDP
- Not clear what polynomial would be reasonable




The non-polynomial case

- Percentage of urban population as a function of per capita GDP
- Not clear what polynomial would be reasonable
- Take log of GDP
- Regression we are computing is
$$y = \theta_0 + \theta_1 \log x_1$$




The non-polynomial case

- Reverse the relationship
- Plot per capita GDP in terms of percentage of urbanization
- Now we take log of the output variable
 $\log y = \theta_0 + \theta_1 x_1$
- Log-linear transformation
- Earlier was linear-log
- Can also use log-log



Regression for classification

- Regression line
- Set a threshold
- Classifier
 - Output below threshold : 0 (No)
 - Output above threshold : 1 (Yes)
- Classifier output is a step function



Smoothen the step


- Sigmoid function

$$\sigma(z) = \frac{1}{1 + e^{-z}}$$

- Input z is output of our regression


$$\sigma(z) = \frac{1}{1 + e^{-(\theta_0 + \theta_1 x_1 + \dots + \theta_k x_k)}}$$

- Adjust parameters to fix horizontal position and steepness of step



Logistic regression

- Compute the coefficients?
- Solve by gradient descent
- Need derivatives to exist
 - Hence smooth sigmoid, not step function
 - $\sigma'(z) = \sigma(z)(1 - \sigma(z))$
- Need a cost function to minimize



Loss function for logistic regression

- Goal is to maximize log likelihood
- Let $h_\theta(x_i) = \sigma(z_i)$. So, $P(y_i = 1 \mid x_i; \theta) = h_\theta(x_i)$,
 $P(y_i = 0 \mid x_i; \theta) = 1 - h_\theta(x_i)$
- Combine as $P(y_i \mid x_i; \theta) = h_\theta(x_i)^{y_i} \cdot (1 - h_\theta(x_i))^{1-y_i}$
- Likelihood: $\mathcal{L}(\theta) = \prod_{i=1}^n h_\theta(x_i)^{y_i} \cdot (1 - h_\theta(x_i))^{1-y_i}$
- Log-likelihood: $\ell(\theta) = \sum_{i=1}^n y_i \log h_\theta(x_i) + (1 - y_i) \log(1 - h_\theta(x_i))$
- Minimize **cross entropy**: $-\sum_{i=1}^n y_i \log h_\theta(x_i) + (1 - y_i) \log(1 - h_\theta(x_i))$

MSE for logistic regression and gradient descent

- Suppose we take mean sum-squared error as the loss function.
- Consider two inputs $x = (x_1, x_2)$

$$C = \frac{1}{n} \sum_{i=1}^n (y_i - \sigma(z_i))^2, \text{ where } z_i = \theta_0 + \theta_1 x_{i1} + \theta_2 x_{i2}$$

- For gradient descent, we compute $\frac{\partial C}{\partial \theta_1}, \frac{\partial C}{\partial \theta_2}, \frac{\partial C}{\partial \theta_0}$


- For $j = 1, 2,$

$$\begin{aligned} \frac{\partial C}{\partial \theta_j} &= \frac{2}{n} \sum_{i=1}^n (y_i - \sigma(z_i)) \cdot -\frac{\partial \sigma(z_i)}{\partial \theta_j} = \frac{2}{n} \sum_{i=1}^n (\sigma(z_i) - y_i) \frac{\partial \sigma(z_i)}{\partial z_i} \frac{\partial z_i}{\partial \theta_j} \\ &= \frac{2}{n} \sum_{i=1}^n (\sigma(z_i) - y_i) \sigma'(z_i) x_{ij} \end{aligned}$$

- $\frac{\partial C}{\partial \theta_0} = \frac{2}{n} \sum_{i=1}^n (\sigma(z_i) - y_i) \frac{\partial \sigma(z_i)}{\partial z_i} \frac{\partial z_i}{\partial b} = \frac{2}{n} \sum_{i=1}^n (\sigma(z_i) - y_i) \sigma'(z_i)$

MSE for logistic regression and gradient descent ...

- For $j = 1, 2$, $\frac{\partial C}{\partial \theta_j} = \frac{2}{n} \sum_{i=1}^n (\sigma(z_i) - y_i) \sigma'(z_i) x_j^i$, and $\frac{\partial C}{\partial \theta_0} = \frac{2}{n} \sum_{i=1}^n (\sigma(z_i) - y_i) \sigma'(z_i)$
- Each term in $\frac{\partial C}{\partial \theta_1}$, $\frac{\partial C}{\partial \theta_2}$, $\frac{\partial C}{\partial \theta_0}$ is proportional to $\sigma'(z_i)$
- Ideally, gradient descent should take large steps when $\sigma(z) - y$ is large
- $\sigma(z)$ is flat at both extremes
- If $\sigma(z)$ is completely wrong, $\sigma(z) \approx (1 - y)$, we still have $\sigma'(z) \approx 0$
- Learning is slow even when current model is far from optimal



Cross entropy and gradient descent

- $C = -[y \ln(\sigma(z)) + (1 - y) \ln(1 - \sigma(z))]$

- $$\begin{aligned}\frac{\partial C}{\partial \theta_j} &= \frac{\partial C}{\partial \sigma} \frac{\partial \sigma}{\partial \theta_j} = - \left[\frac{y}{\sigma(z)} - \frac{1 - y}{1 - \sigma(z)} \right] \frac{\partial \sigma}{\partial \theta_j} \\ &= - \left[\frac{y}{\sigma(z)} - \frac{1 - y}{1 - \sigma(z)} \right] \frac{\partial \sigma}{\partial z} \frac{\partial z}{\partial \theta_j} \\ &= - \left[\frac{y}{\sigma(z)} - \frac{1 - y}{1 - \sigma(z)} \right] \sigma'(z) x_j \\ &= - \left[\frac{y(1 - \sigma(z)) - (1 - y)\sigma(z)}{\sigma(z)(1 - \sigma(z))} \right] \sigma'(z) x_j\end{aligned}$$

Cross entropy and gradient descent ...

- $\frac{\partial C}{\partial \theta_j} = - \left[\frac{y(1 - \sigma(z)) - (1 - y)\sigma(z)}{\sigma(z)(1 - \sigma(z))} \right] \sigma'(z)x_j$
- Recall that $\sigma'(z) = \sigma(z)(1 - \sigma(z))$
- Therefore,
$$\begin{aligned}\frac{\partial C}{\partial \theta_j} &= -[y(1 - \sigma(z)) - (1 - y)\sigma(z)]x_j \\ &= -[y - y\sigma(z) - \sigma(z) + y\sigma(z)]x_j \\ &= (\sigma(z) - y)x_j\end{aligned}$$
- Similarly, $\frac{\partial C}{\partial \theta_0} = (\sigma(z) - y)$
- Thus, as we wanted, the gradient is proportional to $\sigma(z) - y$
- The greater the error, the faster the learning rate