

Chapter 10 – Introduction to Artificial Neural Networks with Keras

This notebook contains all the sample code and solutions to the exercises in chapter 10.



Run in Google Colab (https://colab.research.google.com/github/ageron/handson-ml2/blob/master/10_neural_nets_with_keras.ipynb)

Setup

First, let's import a few common modules, ensure Matplotlib plots figures inline and prepare a function to save the figures. We also check that Python 3.5 or later is installed (although Python 2.x may work, it is deprecated so we strongly recommend you use Python 3 instead), as well as Scikit-Learn ≥ 0.20 and TensorFlow ≥ 2.0 .

```
In [1]: # Python  $\geq 3.5$  is required
import sys
assert sys.version_info >= (3, 5)

# Scikit-Learn  $\geq 0.20$  is required
import sklearn
assert sklearn.__version__ >= "0.20"

try:
    # %tensorflow_version only exists in Colab.
    %tensorflow_version 2.x
except Exception:
    pass

# TensorFlow  $\geq 2.0$  is required
import tensorflow as tf
assert tf.__version__ >= "2.0"

# Common imports
import numpy as np
import os

# to make this notebook's output stable across runs
np.random.seed(42)

# To plot pretty figures
%matplotlib inline
import matplotlib as mpl
import matplotlib.pyplot as plt
mpl.rc('axes', labelsize=14)
mpl.rc('xtick', labelsize=12)
mpl.rc('ytick', labelsize=12)

# Where to save the figures
PROJECT_ROOT_DIR = "."
CHAPTER_ID = "ann"
IMAGES_PATH = os.path.join(PROJECT_ROOT_DIR, "images", CHAPTER_ID)
os.makedirs(IMAGES_PATH, exist_ok=True)

def save_fig(fig_id, tight_layout=True, fig_extension="png", resolution=300):
    path = os.path.join(IMAGES_PATH, fig_id + "." + fig_extension)
    print("Saving figure", fig_id)
    if tight_layout:
        plt.tight_layout()
    plt.savefig(path, format=fig_extension, dpi=resolution)
```

```
# Ignore useless warnings (see SciPy issue #5998)
import warnings
warnings.filterwarnings(action="ignore", message="^internal gelsd")
```

Perceptrons

Note: we set `max_iter` and `tol` explicitly to avoid warnings about the fact that their default value will change in future versions of Scikit-Learn.

```
In [2]: import numpy as np
from sklearn.datasets import load_iris
from sklearn.linear_model import Perceptron

iris = load_iris()
X = iris.data[:, (2, 3)] # petal length, petal width
y = (iris.target == 0).astype(np.int)

per_clf = Perceptron(max_iter=1000, tol=1e-3, random_state=42)
per_clf.fit(X, y)

y_pred = per_clf.predict([[2, 0.5]])

/home/madhavan/miniconda3/lib/python3.7/site-packages/ipykernel_launcher.py:7:
DeprecationWarning: `np.int` is a deprecated alias for the builtin `int`. To s
ilence this warning, use `int` by itself. Doing this will not modify any behav
ior and is safe. When replacing `np.int`, you may wish to use e.g. `np.int64`
or `np.int32` to specify the precision. If you wish to review your current us
e, check the release note link for additional information.
Deprecated in NumPy 1.20; for more details and guidance: https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations (https://numpy.org/devdocs/release/1.20.0-notes.html#deprecations)
import sys
```

```
In [3]: y_pred
```

```
Out[3]: array([1])
```

```
In [4]: a = -per_clf.coef_[0][0] / per_clf.coef_[0][1]
b = -per_clf.intercept_ / per_clf.coef_[0][1]

axes = [0, 5, 0, 2]


x0, x1 = np.meshgrid(
    np.linspace(axes[0], axes[1], 500).reshape(-1, 1),
    np.linspace(axes[2], axes[3], 200).reshape(-1, 1),
)
X_new = np.c_[x0.ravel(), x1.ravel()]
y_predict = per_clf.predict(X_new)
zz = y_predict.reshape(x0.shape)

plt.figure(figsize=(10, 4))
plt.plot(X[y==0, 0], X[y==0, 1], "bs", label="Not Iris-Setosa")
plt.plot(X[y==1, 0], X[y==1, 1], "yo", label="Iris-Setosa")

plt.plot([axes[0], axes[1]], [a * axes[0] + b, a * axes[1] + b], "k-", linewidth=2)
from matplotlib.colors import ListedColormap
custom_cmap = ListedColormap(['#9898ff', '#fafab0'])

plt.contourf(x0, x1, zz, cmap=custom_cmap)
plt.xlabel("Petal length", fontsize=14)
plt.ylabel("Petal width", fontsize=14)
plt.legend(loc="lower right", fontsize=14)
plt.axis(axes)
```

```
save_fig("perceptron_iris_plot")  
plt.show()  
Saving figure perceptron_iris_plot
```



Activation functions

```
In [5]: def sigmoid(z):  
        return 1 / (1 + np.exp(-z))  
  
        def relu(z):  
            return np.maximum(0, z)  
  
        def derivative(f, z, eps=0.000001):  
            return (f(z + eps) - f(z - eps)) / (2 * eps)
```

```

In [6]: z = np.linspace(-5, 5, 200)

plt.figure(figsize=(11,4))


plt.subplot(121)
plt.plot(z, np.sign(z), "r-", linewidth=1, label="Step")
plt.plot(z, sigmoid(z), "g--", linewidth=2, label="Sigmoid")
plt.plot(z, np.tanh(z), "b-", linewidth=2, label="Tanh")
plt.plot(z, relu(z), "m-.", linewidth=2, label="ReLU")
plt.grid(True)
plt.legend(loc="center right", fontsize=14)
plt.title("Activation functions", fontsize=14)
plt.axis([-5, 5, -1.2, 1.2])

plt.subplot(122)
plt.plot(z, derivative(np.sign, z), "r-", linewidth=1, label="Step")
plt.plot(0, 0, "ro", markersize=5)
plt.plot(0, 0, "rx", markersize=10)
plt.plot(z, derivative(sigmoid, z), "g--", linewidth=2, label="Sigmoid")
plt.plot(z, derivative(np.tanh, z), "b-", linewidth=2, label="Tanh")
plt.plot(z, derivative(relu, z), "m-.", linewidth=2, label="ReLU")
plt.grid(True)
#plt.legend(loc="center right", fontsize=14)
plt.title("Derivatives", fontsize=14)
plt.axis([-5, 5, -0.2, 1.2])

save_fig("activation_functions_plot")
plt.show()

```

Saving figure activation_functions_plot



Building an Image Classifier

First let's import TensorFlow and Keras.

```

In [9]: import tensorflow as tf
        from tensorflow import keras

```

```

In [10]: tf. version

```

```

Out[10]: '2.4.1'

```

```

In [11]: keras. version

```

```

Out[11]: '2.4.0'

```

Let's start by loading the fashion MNIST dataset. Keras has a number of functions to load popular datasets in `keras.datasets`. The dataset is already split for you between a training set and a test set, but it can be useful to split the training set further to have a validation set:

```
In [12]: fashion_mnist = keras.datasets.fashion_mnist
(X_train_full, y_train_full), (X_test, y_test) = fashion_mnist.load_data()
```

The training set contains 60,000 grayscale images, each 28x28 pixels:

```
In [13]: X_train_full.shape
```

```
Out[13]: (60000, 28, 28)
```

Each pixel intensity is represented as a byte (0 to 255):

```
In [14]: X_train_full.dtype
```

```
Out[14]: dtype('uint8')
```

Let's split the full training set into a validation set and a (smaller) training set. We also scale the pixel intensities down to the 0-1 range and convert them to floats, by dividing by 255.

```
In [15]: X_valid, X_train = X_train_full[:5000] / 255., X_train_full[5000:] / 255.
y_valid, y_train = y_train_full[:5000], y_train_full[5000:]
X_test = X_test / 255.
```

You can plot an image using Matplotlib's `imshow()` function, with a `'binary'` color map:

```
In [16]: plt.imshow(X_train[0], cmap="binary")
plt.axis('off')
plt.show()
```



The labels are the class IDs (represented as uint8), from 0 to 9:

```
In [17]: y_train
```

```
Out[17]: array([4, 0, 7, ..., 3, 0, 5], dtype=uint8)
```

Here are the corresponding class names:

```
In [18]: class_names = ["T-shirt/top", "Trouser", "Pullover", "Dress", "Coat",
                        "Sandal", "Shirt", "Sneaker", "Bag", "Ankle boot"]
```

So the first image in the training set is a coat:

```
In [19]: class_names[y_train[0]]
```

```
Out[19]: 'Coat'
```

The validation set contains 5,000 images, and the test set contains 10,000 images:

```
In [20]: X_valid.shape
```

```
Out[20]: (5000, 28, 28)
```

```
In [21]: X_test.shape
```

```
Out[21]: (10000, 28, 28)
```

Let's take a look at a sample of the images in the dataset:

```
In [22]: n_rows = 4
n_cols = 10
plt.figure(figsize=(n_cols * 1.2, n_rows * 1.2))
for row in range(n_rows):
    for col in range(n_cols):
        index = n_cols * row + col
        plt.subplot(n_rows, n_cols, index + 1)
        plt.imshow(X_train[index], cmap="binary", interpolation="nearest")
        plt.axis('off')
        plt.title(class_names[y_train[index]], fontsize=12)
plt.subplots_adjust(wspace=0.2, hspace=0.5)
save_fig('fashion_mnist_plot', tight_layout=False)
plt.show()
```

Saving figure fashion_mnist_plot



```
In [23]: model = keras.models.Sequential()
model.add(keras.layers.Flatten(input_shape=[28, 28]))
model.add(keras.layers.Dense(300, activation="relu"))
model.add(keras.layers.Dense(100, activation="relu"))
model.add(keras.layers.Dense(10, activation="softmax"))
```

```
In [24]: keras.backend.clear_session()
np.random.seed(42)
tf.random.set_seed(42)
```

```
In [25]: model = keras.models.Sequential([
    keras.layers.Flatten(input_shape=[28, 28]),
    keras.layers.Dense(300, activation="relu"),
    keras.layers.Dense(100, activation="relu"),
    keras.layers.Dense(10, activation="softmax")
])
```

```
In [26]: model.layers
```

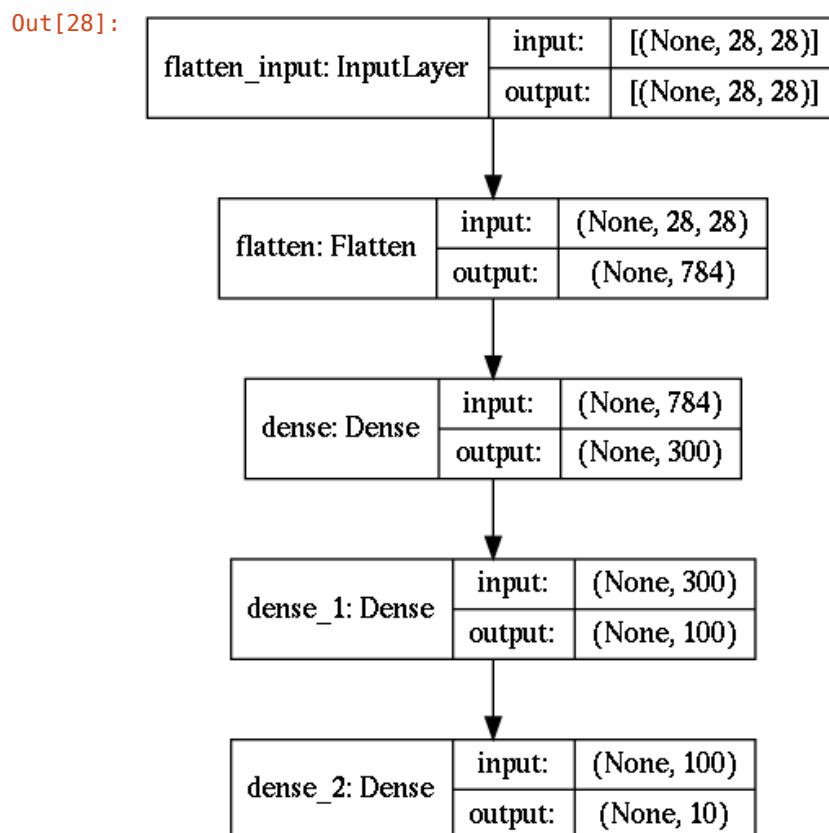
```
Out[26]: [<tensorflow.python.keras.layers.core.Flatten at 0x7f15546140d0>,
<tensorflow.python.keras.layers.core.Dense at 0x7f1554614350>,
<tensorflow.python.keras.layers.core.Dense at 0x7f1554614690>,
<tensorflow.python.keras.layers.core.Dense at 0x7f155506a310>]
```

```
In [27]: model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 784)	0
dense (Dense)	(None, 300)	235500
dense_1 (Dense)	(None, 100)	30100
dense_2 (Dense)	(None, 10)	1010
Total params: 266,610		
Trainable params: 266,610		
Non-trainable params: 0		

```
In [28]: keras.utils.plot_model(model, "my_mnist_model.png", show_shapes=True)
```




```
model.compile(loss=keras.losses.sparse_categorical_crossentropy,
```

```
In [37]: history = model.fit(X_train, y_train, epochs=30,  
                             validation data=(X valid, y valid))
```

```

Epoch 1/30
1719/1719 [=====] - 17s 10ms/step - loss: 1.0188 - ac
curacy: 0.6805 - val_loss: 0.5218 - val_accuracy: 0.8210
Epoch 2/30
1719/1719 [=====] - 11s 6ms/step - loss: 0.5028 - acc
uracy: 0.8259 - val_loss: 0.4351 - val_accuracy: 0.8530
Epoch 3/30
1719/1719 [=====] - 15s 9ms/step - loss: 0.4483 - acc
uracy: 0.8423 - val_loss: 0.5354 - val_accuracy: 0.7994
Epoch 4/30
1719/1719 [=====] - 17s 10ms/step - loss: 0.4207 - ac
curacy: 0.8525 - val_loss: 0.3917 - val_accuracy: 0.8656
Epoch 5/30
1719/1719 [=====] - 16s 9ms/step - loss: 0.4059 - acc
uracy: 0.8587 - val_loss: 0.3748 - val_accuracy: 0.8688
Epoch 6/30
1719/1719 [=====] - 12s 7ms/step - loss: 0.3753 - acc
uracy: 0.8678 - val_loss: 0.3711 - val_accuracy: 0.8724
Epoch 7/30
1719/1719 [=====] - 14s 8ms/step - loss: 0.3651 - acc
uracy: 0.8710 - val_loss: 0.3629 - val_accuracy: 0.8732
Epoch 8/30
1719/1719 [=====] - 14s 8ms/step - loss: 0.3477 - acc
uracy: 0.8753 - val_loss: 0.3837 - val_accuracy: 0.8626
Epoch 9/30
1719/1719 [=====] - 19s 11ms/step - loss: 0.3480 - ac
curacy: 0.8760 - val_loss: 0.3596 - val_accuracy: 0.8698
Epoch 10/30

```

In [38]: `history.params`

Out[38]: `{'verbose': 1, 'epochs': 30, 'steps': 1719}`

In [39]: `print(history.epoch)`

```
[0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19, 20, 21,
22, 23, 24, 25, 26, 27, 28, 29]
```


In [40]: `history.history.keys()`

Out[40]: `dict_keys(['loss', 'accuracy', 'val_loss', 'val_accuracy'])`

In [41]: `import pandas as pd`

```
pd.DataFrame(history.history).plot(figsize=(8, 5))
plt.grid(True)
plt.gca().set_ylim(0, 1)
save_fig("keras_learning_curves_plot")
plt.show()
```

Saving figure keras_learning_curves_plot



In [42]: `model.evaluate(X_test, y_test)`

313/313 [=====] - 1s 4ms/step - loss: 0.3348 - accuracy: 0.8831

Out[42]: [0.33482903242111206, 0.8830999732017517]

In [43]: `X_new = X_test[:3]`
`y_proba = model.predict(X_new)`
`y_proba.round(2)`

Out[43]: array([[0. , 0. , 0. , 0. , 0. , 0.01, 0. , 0.02, 0. , 0.96],
 [0. , 0. , 0.99, 0. , 0.01, 0. , 0. , 0. , 0. , 0.],
 [0. , 1. , 0. , 0. , 0. , 0. , 0. , 0. , 0. , 0.]],
 dtype=float32)

In [44]: `y_pred = model.predict_classes(X_new)`
`y_pred`

/home/madhavan/miniconda3/lib/python3.7/site-packages/tensorflow/python/keras/engine/sequential.py:450: UserWarning: `model.predict_classes()` is deprecated and will be removed after 2021-01-01. Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your model does multi-class classification (e.g. if it uses a `softmax` last-layer activation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary classification (e.g. if it uses a `sigmoid` last-layer activation).
 warnings.warn("`model.predict_classes()` is deprecated and "

Out[44]: array([9, 2, 1])

In [45]: `np.array(class_names)[y_pred]`

Out[45]: array(['Ankle boot', 'Pullover', 'Trouser'], dtype='<U11')

```
In [46]: y_new = y_test[:3]
y_new
```

```
Out[46]: array([9, 2, 1], dtype=uint8)
```

```
In [47]: plt.figure(figsize=(7.2, 2.4))
for index, image in enumerate(X_new):
    plt.subplot(1, 3, index + 1)
    plt.imshow(image, cmap="binary", interpolation="nearest")
    plt.axis('off')
    plt.title(class_names[y_test[index]], fontsize=12)
plt.subplots_adjust(wspace=0.2, hspace=0.5)
save_fig('fashion_mnist_images_plot', tight_layout=False)
plt.show()
```

Saving figure fashion_mnist_images_plot

